

Developing an Extensible, Portable, Scalable Toolkit for Massively Parallel Incompressible Smoothed Particle Hydrodynamics (ISPH)

Dr Xiaohu Guo
Hartree Centre, STFC

Dr Benedict Rogers, Dr Steven Lind, Prof. Peter Stansby
SPH Group, School of Mechanical, Aerospace and Civil
Engineering, University of Manchester

Overview

- Motivations
- Data preconditioning for extensibility and neighbour list searching
- Domain Decomposition and Dynamic load balancing with unstructured communication toolkit
- Particles reordering for ISPH
- Scalable sparse Linear solver for ISPH
- Applications and Summary

- The stability, accuracy, energy conservation, boundary conditions of ISPH have been greatly improved (*S. Lind, H. Gotoh and A. Khayyer*)
- Computational challenges:
 - Particles irregular distribution due to solving complex, nonlinear and distorted flow.
 - Inherited SPH intrinsic computing bounded property.
 - Due to requirement solving pressure Poisson equation using sparse linear solver, ISPH is both memory bounded and computing bounded.

A Toolkit to facilitate ISPH software development for scientific users.

Incompressible SPH Implementation

- ISPH_DF – A Divergence-free Projection method in SPH

$$\mathbf{r}_i^* = \mathbf{r}_i^n + \delta t \mathbf{u}_i^n$$

Particles are advected to \mathbf{r}_i^*

LOOP 1

$$\mathbf{u}_i^* = \mathbf{u}_i^n + \left(\nu \nabla^2 \mathbf{u}_i^n + \mathbf{F}_i^n \right) \delta t$$

Calculate \mathbf{u}_i^* without pressure gradient.

LOOP 2

$$\nabla \cdot \left(\frac{1}{\rho} \nabla p^{n+1} \right)_i = \frac{1}{\delta t} \nabla \cdot \mathbf{u}_i^*$$

Solve pressure Poisson equation to get pressure P_i^{n+1} .

LOOP 3

$$\mathbf{u}_i^{n+1} = \mathbf{u}_i^* - \frac{\delta t}{\rho} \nabla p_i^{n+1}$$

Calculate velocity \mathbf{u}_i^{n+1} .

$$\mathbf{r}_i^{n+1} = \mathbf{r}_i^n + \delta t \left(\frac{\mathbf{u}_i^{n+1} + \mathbf{u}_i^n}{2} \right)$$

Particles' positions are centered in time.

LOOP 4

$$\delta \mathbf{r}_s = -\mathcal{D} \left(\frac{\partial \mathcal{C}}{\partial s} \mathbf{s} + \alpha \left(\frac{\partial \mathcal{C}}{\partial n} - \beta \right) \mathbf{n} \right)$$

Shift particles to maintain stability.

*ISPH (Cummins & Rudman 1999)

● Fluid ● Boundary ● Rigid Body

offset

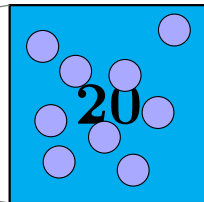
cell vec

count

Provide extensibility if multiple types of particles involved.

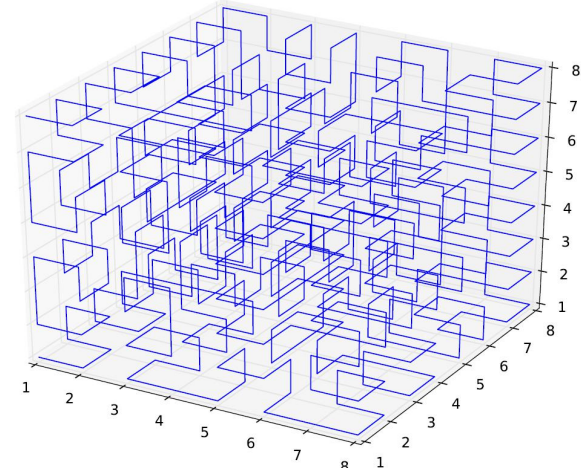
Domain Decomposition and Dynamic load balancing with SFC

57	58	59	60	61	62	63	64
49	50	51	52	53	54	55	56
41	42	43	44	45	46	47	48
33	34	35	36	37	38	39	40
25	26	27	28	29	30	31	32
17	18	19	20	21	22	23	24
9	10	11	12	13	14	15	16
1	2	3	4	5	6	7	8

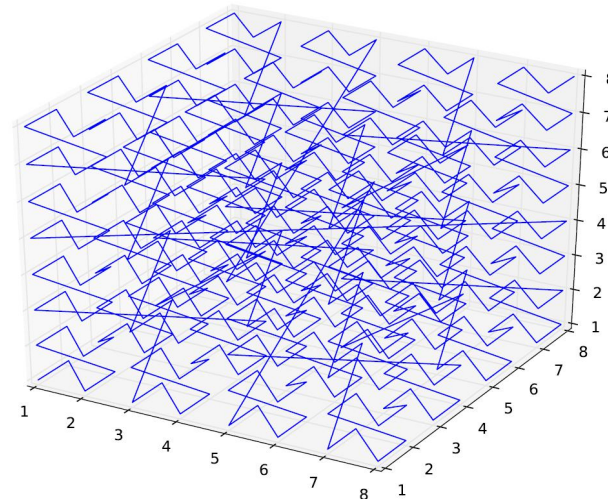


Partitions: P_0 =blue, P_1 =green, P_2 =yellow, P_3 =pink

We use the cell as partition object, calculate SFC according to cells

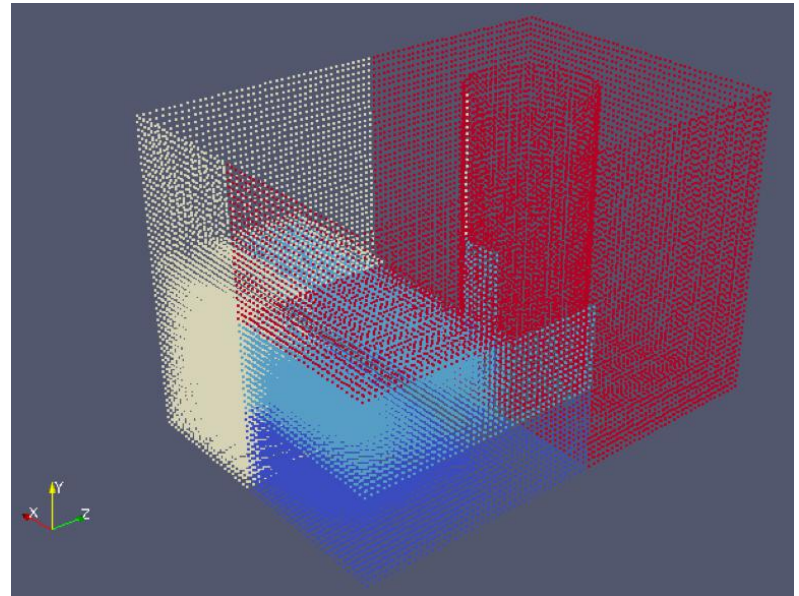
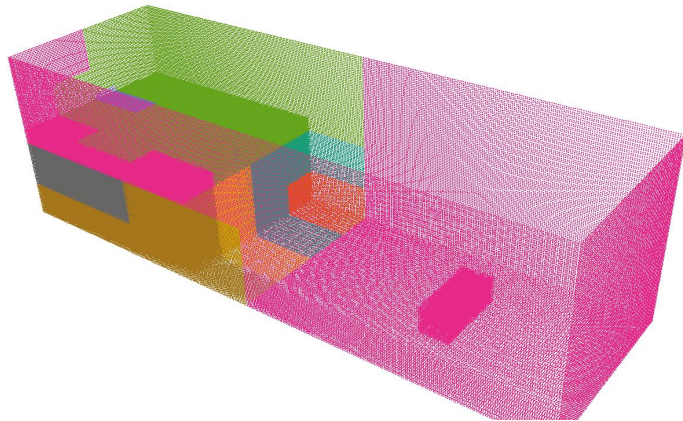


Hilbert Space Filling Curve:



Morton Space Filling Curve:

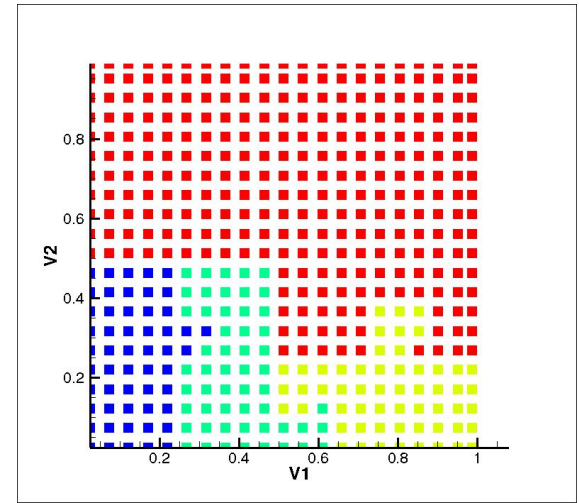
There is no extra effort for the domain decomposition and dynamic load balancing for solid objects



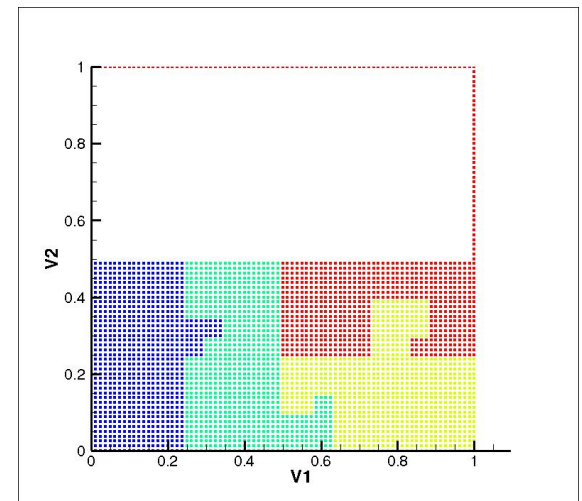
Colours denote partitions

Unstructured communication

- As a result of using space filling curve, the domain shape become irregular.
- Each processor has cells (contain particles) to send to other processors, but no processor knows what messages it will receive.
- ISPH produces a data structure called a **communication plan** which encapsulates the basic information about the communication operation



Cells Decomposition



Particles Decomposition

ISPH Domain Decomposition and Dynamic load balancing software reengineering

Halo Exchange

	34	35	36	37
	26	27	28	29
17	18	19	20	21
9	10	11	12	13
1	2	3	4	5

Halo Send and Receive Plan for cells:

- P_0 :SEND(to send to P_3): 9,10,11,19,27,28
- P_0 :SEND(to send to P_2): 20,28
- P_0 :SEND(to send to P_1): 4,12,20,28
- P_0 :RECV(from P_3): 17,18,26,32,35,36,37
- P_0 :RECV(from P_2): 29
- P_0 :RECV(from P_1): 5,13,21

Partitions: P_0 =blue, P_1 =green, P_2 =yellow, P_3 =pink

Partition P0 halo plan setup

ISPH Halo Exchange initial implementation

New Syntax of ISPH Halo Exchange

Do $t = 1 \rightarrow \text{total_number_of_timesteps}$

!! setup halo exchange plan

call `halo_plan_setup`

`recvs)`

!! update halo objects `integer`

call `halo_update_integer()`

!! update the physical data blocks to

!! an appropriate partition

call `halo_update_real`

perform local calculation

!! destroy halo exchange plan if

!! partition changed

If (change) then

call `halo_plan_destroy`(halo sends, halo

`recvs)`

endif

end do

No Big Deal !

Do $t = 1 \rightarrow \text{total_number_of_timesteps}$

!! setup halo exchange plan

call `halo_plan_setup`(halo_sends,

!! update the physical field data to

!! an appropriate partition

call `halo_update`(halo_sends, halo_recvs,

`Var)`

perform local calculation for Var

call `halo_plan_destroy` if

change plan if

ed

call `halo_plan_destroy`(halo sends, halo

`recvs)`

endif

end do

**Big deal in terms
of new model
development !**

**Two Major steps to reduce memory
footprint.**

Not saving the neighbour list

LOOP 1 $\mathbf{u}_i^* = \mathbf{u}_i^n + \left(\nu \nabla^2 \mathbf{u}_i^n + \mathbf{F}_i^n \right) \delta t$

LOOP 2 $\nabla \cdot \left(\frac{1}{\rho} \nabla p^{n+1} \right)_i = \frac{1}{\delta t} \nabla \cdot \mathbf{u}_i^*$

LOOP 3 $\mathbf{u}_i^{n+1} = \mathbf{u}_i^* - \frac{\delta t}{\rho} \nabla p_i^{n+1}$

LOOP 4 $\delta \mathbf{r}_s = -\mathcal{D} \left(\frac{\partial \mathcal{C}}{\partial \mathbf{s}} \mathbf{s} + \alpha \left(\frac{\partial \mathcal{C}}{\partial \mathbf{n}} - \beta \right) \mathbf{n} \right)$

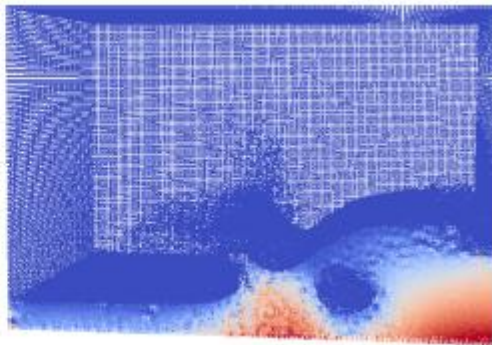
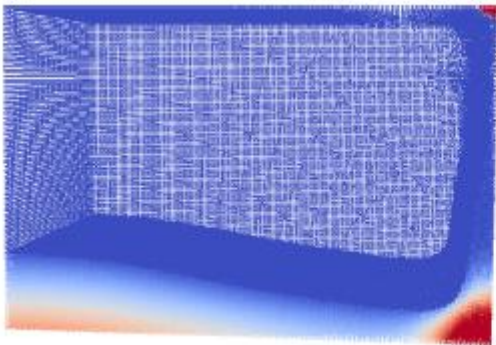
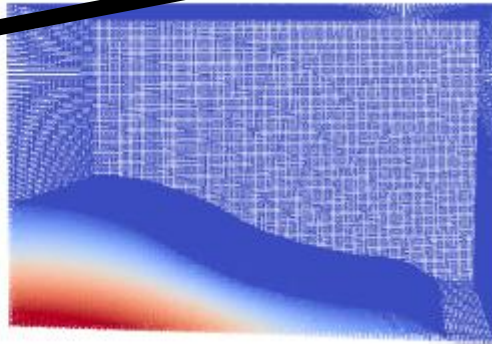
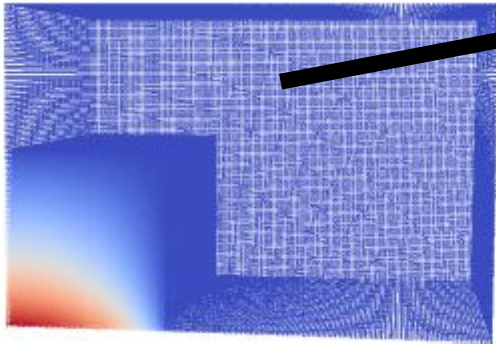
- Instead looping through all particles i , we loop each cell
- In each cell, we then have a subloop to calculate all properties for each particle i
- Such way, we don't need save each particle i 's neighbour list.
- Within each loop, need reduce redundant calculations,

NOT saving neighbour list is much faster and also provide solutions for particles ordering.

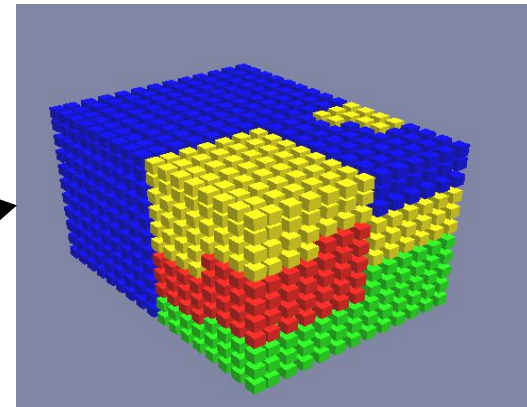
Bounding box creation for MPI partition mapping

The number of cells can be huge.

e.g. for 3D dam break, we have 1 billion particles, the number of cells is $\approx 1/300$ billion

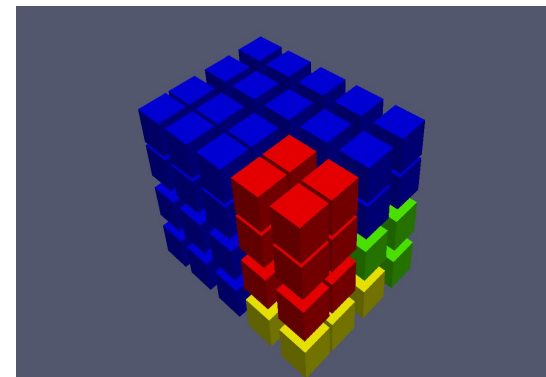


Cells



ble

ells

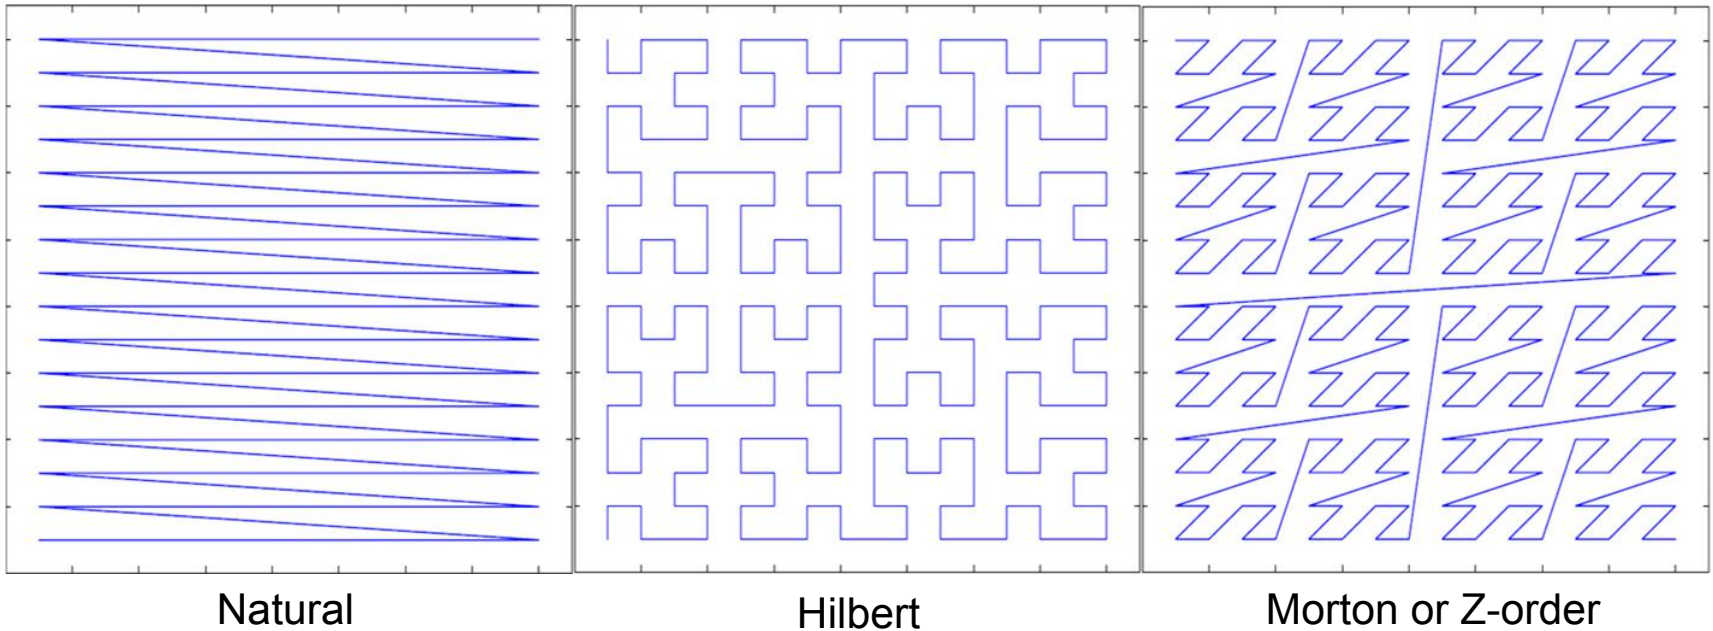


Key Point:

**Number of
Bounding Boxes**

$\approx 20\%$ Number of $2h$ cells

Particles reordering for ISPH



Particle Reordering – Why?

By reordering particles to be close to each other in memory to reduce cache access & hence reduce computation time

How?

Particles order are ordered by ordering the cells that contain them using either Natural, Hilbert SFC or Morton SFC.

→ **GAIN:** Light overhead to calculate SFC using bit shifting

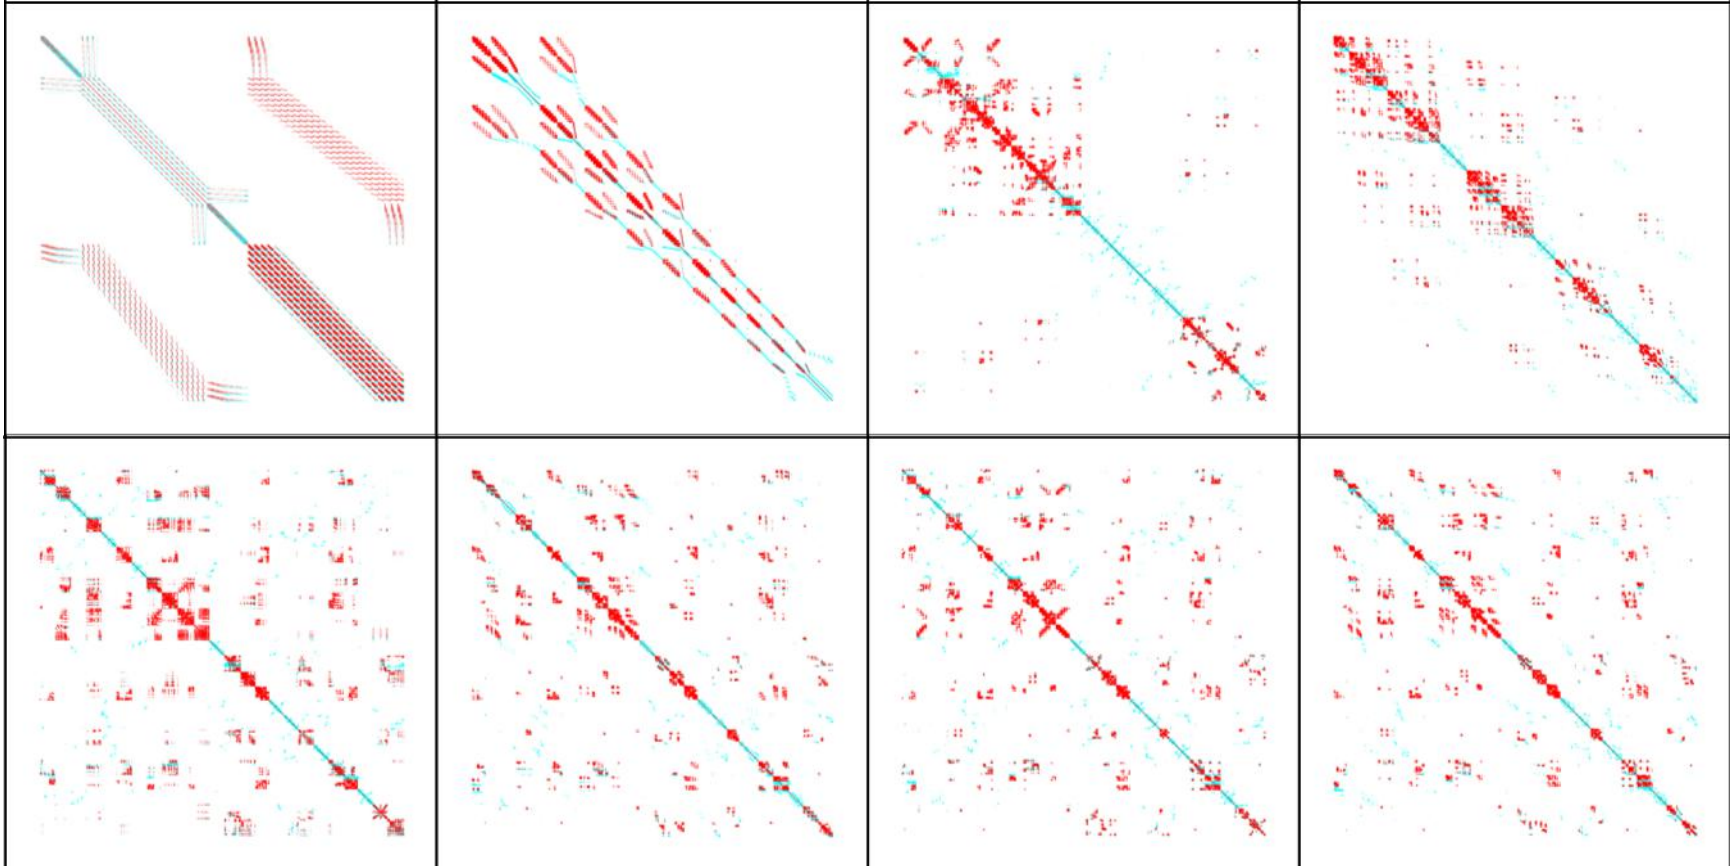
Effects of sorting on 3D PPE matrix structure

No sort

Natural

Hilbert

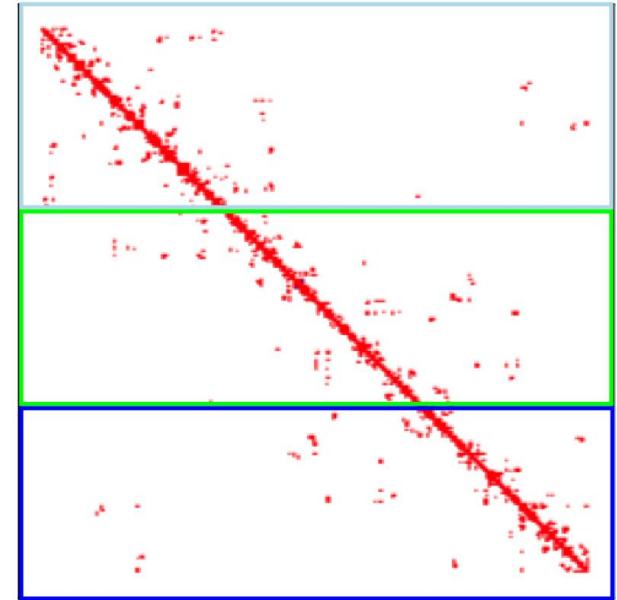
Morton



With particles sorting, in Serial, reduce matrix bandwidth. In parallel, needs efficient global particles sorting—expensive !

Flexible Sparse linear solver implementation using PETSc

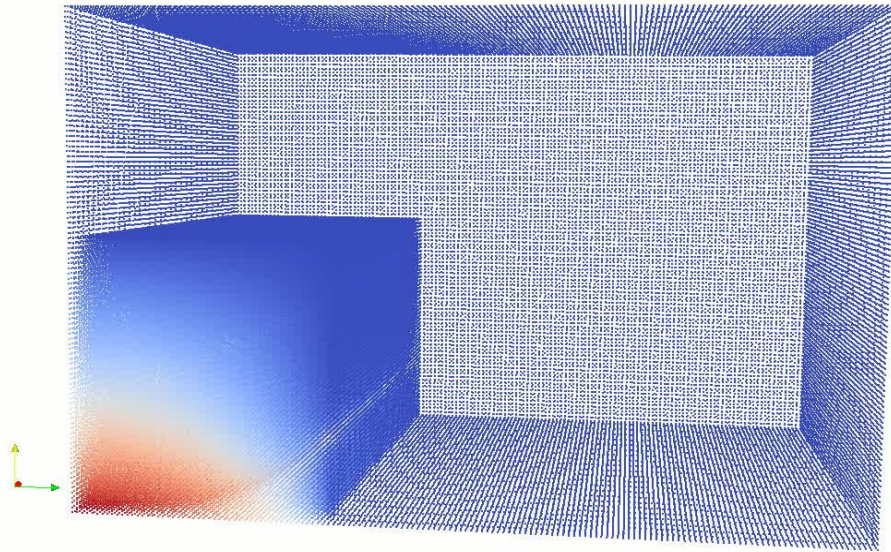
- We use sparse linear solver library PETSc which provide rich options of Krylov subspace solvers and preconditioners.
- Assembly by each row with MatsetValues, not MatsetValue
- Involves set boundary values, correction method, this has be done before MatAssemblyBegin/End, to reduce memory access to PETSc memory space.
- Provide a efficient way to explore PETSc preconditioners and Krylov subspace solvers



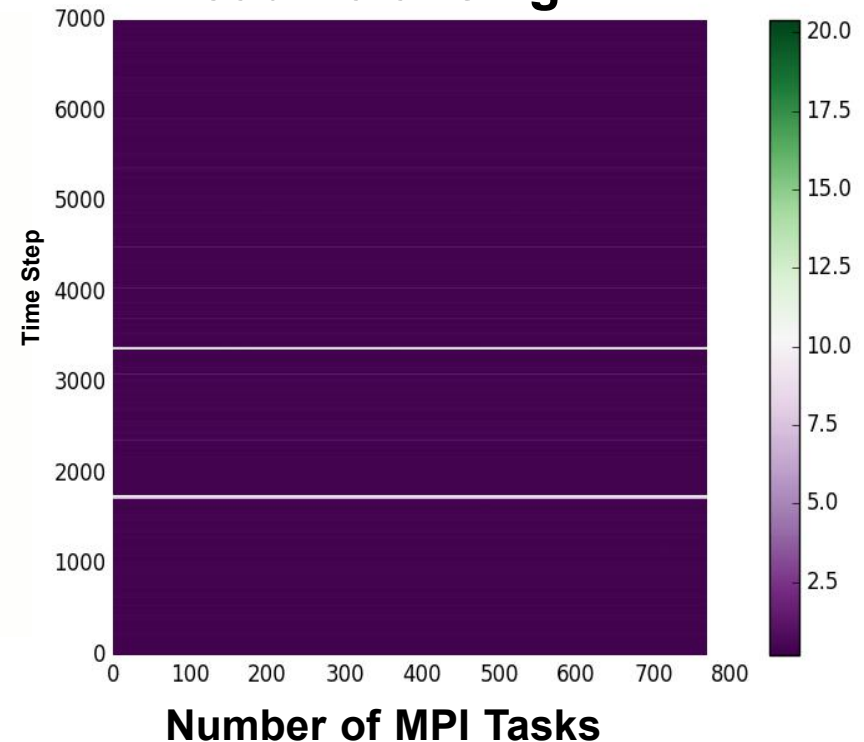
PETSc partition matrix by rows, each colour represent the partition of the PPE matrix

```
mpirun -np 3 isph3d -pc_type hypre  
-pc_hypre_type boomeramg  
-ksp_type gmre
```

Dynamic Load Balancing for 3D Dambreak with Dry Bed using ISPH



**Heat Map to show
Load Balancing**



With around 100 million particles using 768 MPI partitions

ISPH3D Overall Performance Analysis: strong scaling

ISPH3D code runtime performance on ARCHER Cray XC30

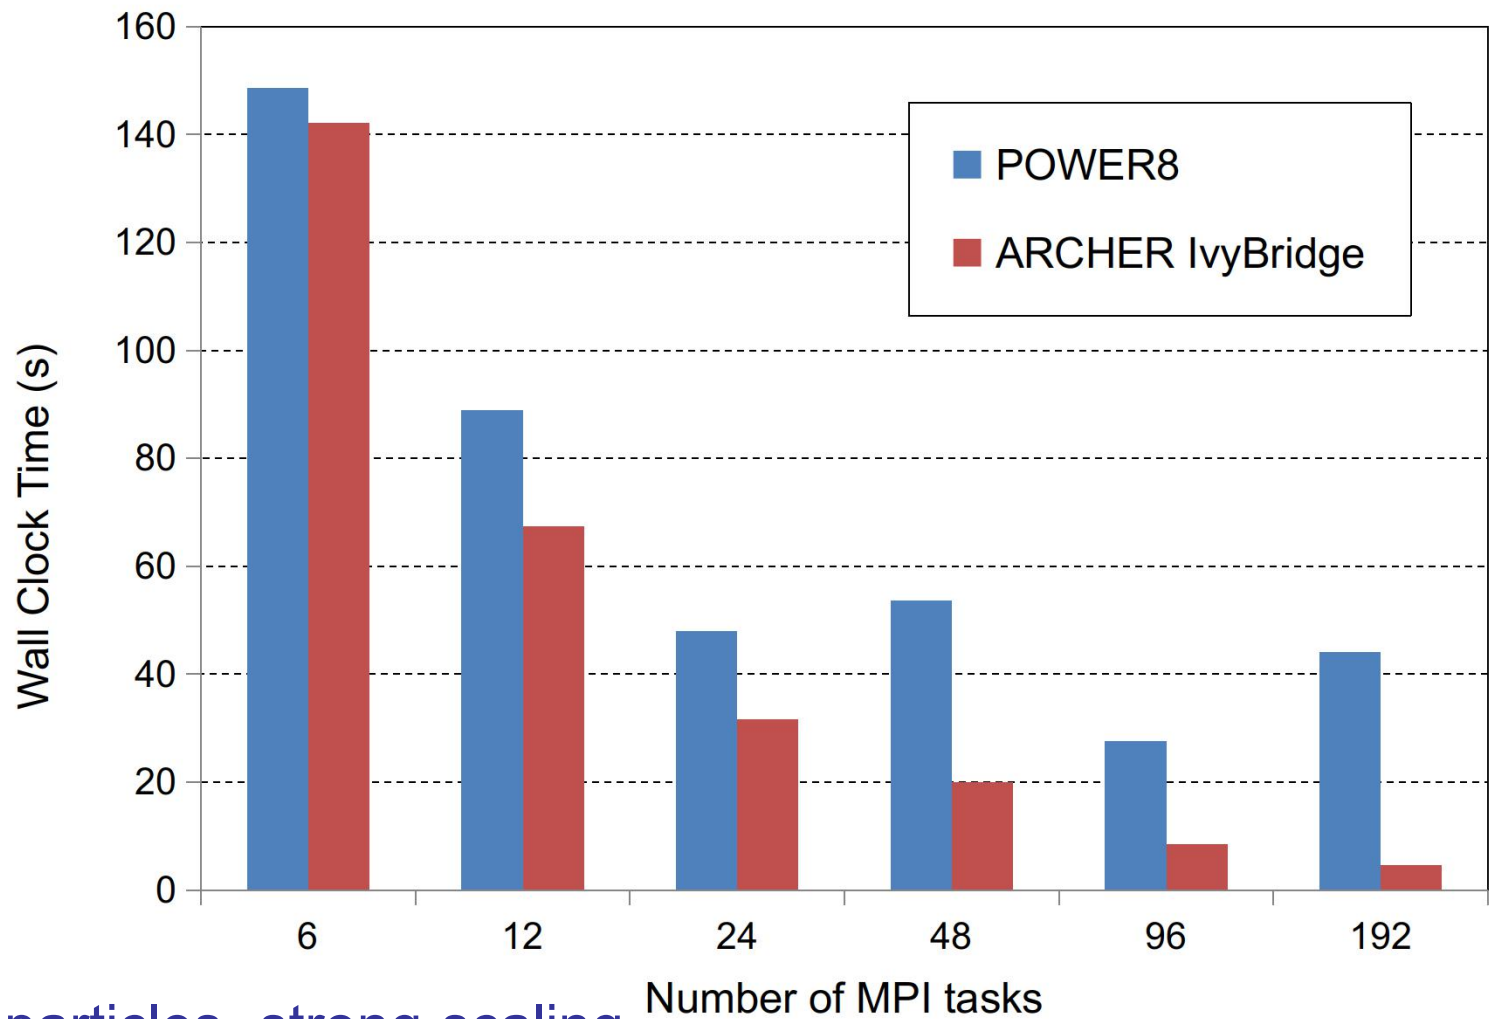
Conclusion

- The ISPH3D can scale well up to 12K cores with strong scaling.
- The SPH related kernels scales almost linearly
- There are now strong needs to improve scaling of PCSETUP, MatAssembly.
- **A fast converging scalable** preconditioner of ISPH

Efficiency



ISPH Performance on IBM Power8

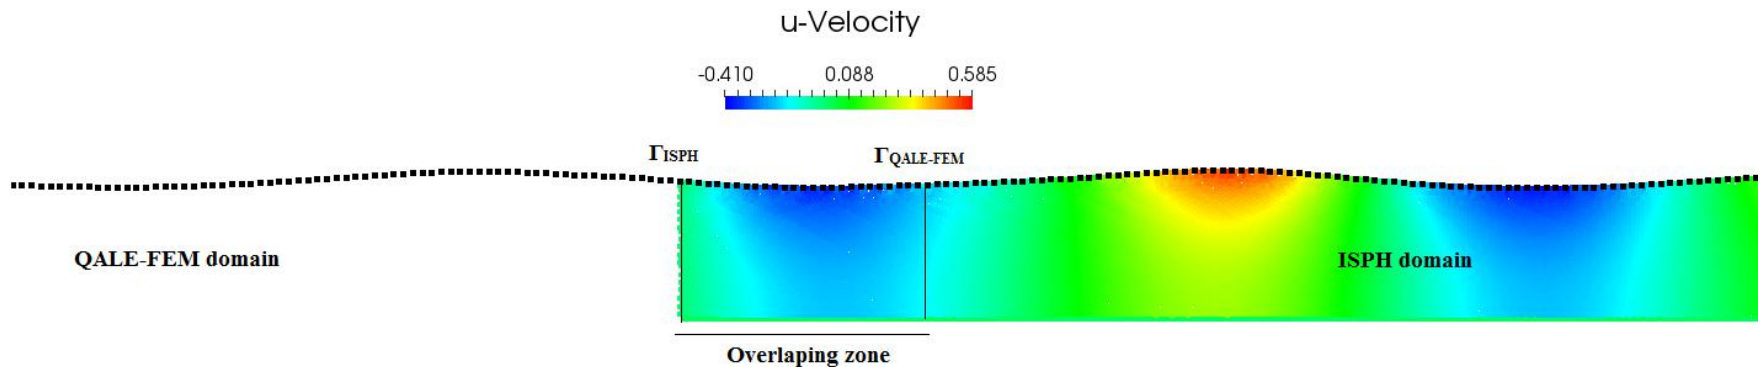
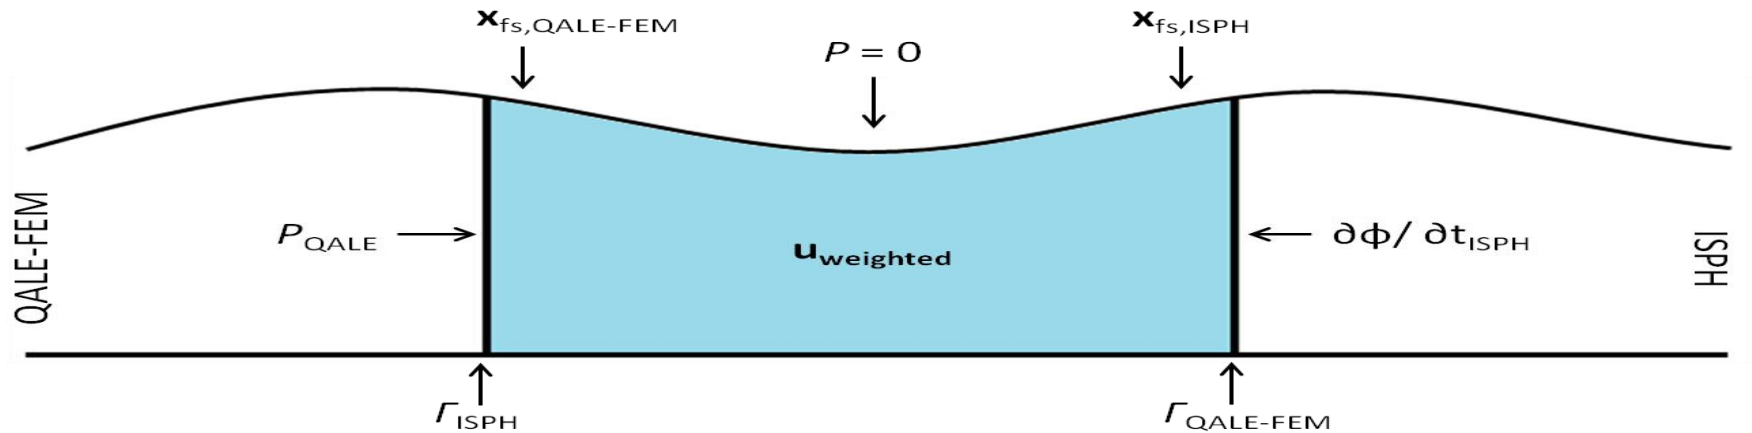


ISPH 5M particles, strong scaling

ARCHER Cray XC30 each node 2x 2.7 GHz, 12-core E5-2697 v2 (Ivy Bridge)

Scales well to 24 cores, SMT helps at 96 virtual cores (SMT=4)

Application: Coupling of ISPH + QALE-FEM



Summary so far ...

- Domain Decomposition and Dynamic load balancing
- Efficient nearest neighbour searching kernel Implementation
- Particles reordering: performance improvements
- Flexible implementation of using PETSc solving PPE for ISPH
- The code can now scale well up to tens of thousands cores which enable the ISPH to work on large-scale real applications
- Although it's implemented together with ISPH, it can also benefit other projection based particle method as a separate parallel library – **separation of concern for MPI based parallelism**

Things happening ...

- **Portability study** for heterogeneous computing architectures.
- Handling **complex geometries**
- Fast converging and scalable preconditioners designed for ISPH
- Boundary conditions, MBT, Inlet/Outlet.
- support of **multiphysics coupling** in large scale.
- **Preprocessing and post processing tools** targeting to complex geometries with large number of particles for ISPH3D
- Efficient parallel **I/O**.
- **Strong scaling** for billion particles above

Acknowledgements

- The authors would also like to acknowledge the funding support under the embedded CSE programme of the ARCHER UK National Supercomputing Service (<http://www.archer.ac.uk>). The authors would also like to thank the EPCC eCSE support team for their help throughout this work.
- The work was partially funded by the EPSRC Grant "SERT: Scale-free, Energy-aware, Resilient and Transparent Adaptation of CSE Applications to Mega-core Systems ". Grant Number:EP/M01147X/1.
- The work was carried out on ISPH software which was original funded by the EPSRC Grant "An incompressible smoothed particle hydrodynamics (ISPH) wave basin with structure interaction for fully nonlinear and extreme coastal waves". Grant Number: EP/H018603/1, EP/H018638/1.



Thanks !